



CECS 326 Operating Systems

Introduction

Hailu Xu

California State University Long Beach

Introduction

- Computer System Structure
- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture
- Operating-System Structure
- Operating-System Operations
- Process Management
- Memory Management
- Storage Management
- Kernel Data Structures
- Computing Environments
- Open-Source Operating Systems

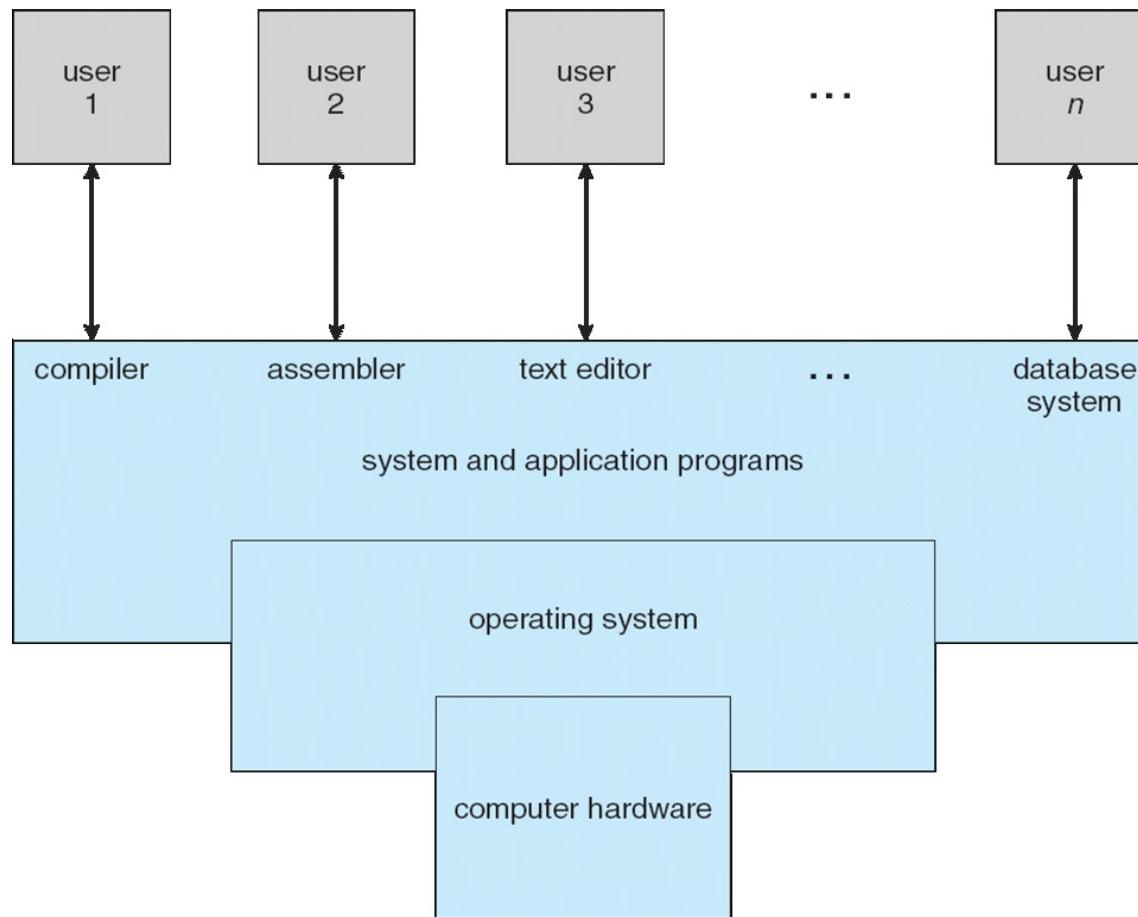
Objectives

- To describe the basic organization of computer systems
- To provide a grand tour of the major components of operating systems
- To give an overview of the many types of computing environments
- To explore several open-source operating systems

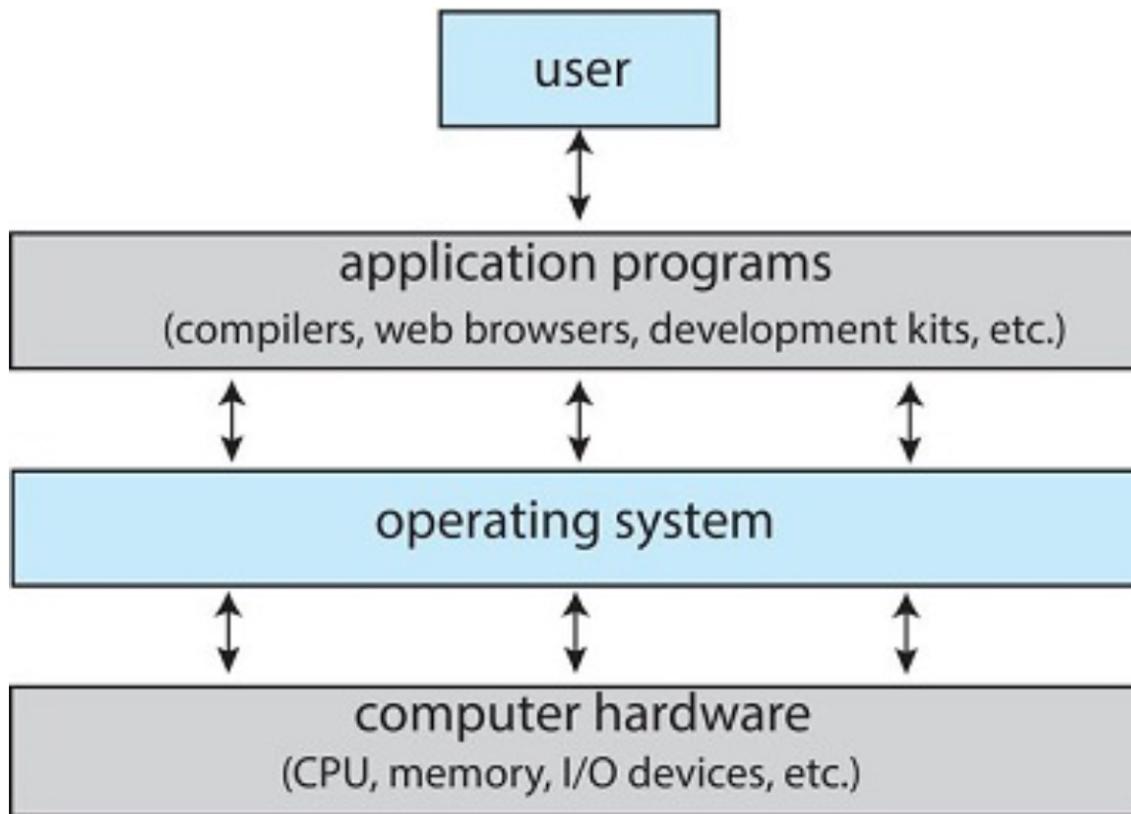
Computer System Structure

- Computer system can be divided into four components:
 - Hardware – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - **Operating system**
 - ▶ Controls and coordinates use of hardware among various applications and users
 - Application programs – define the ways in which the system resources are used to solve the computing problems of the users
 - ▶ Word processors, compilers, web browsers, database systems, video games
 - Users
 - ▶ People, machines, other computers

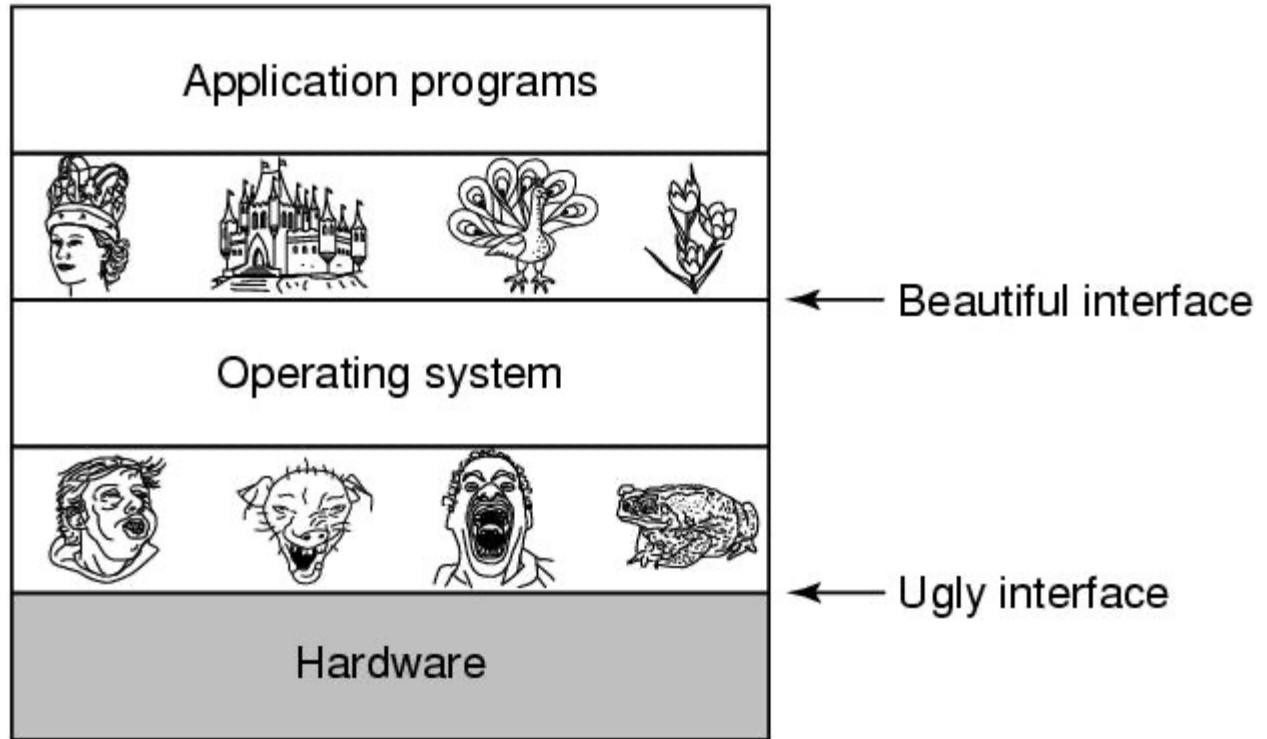
Four Components of a Computer System



Computer System Structure



The Operating System as an Extended Machine



What is an Operating System?

- A **program** that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner

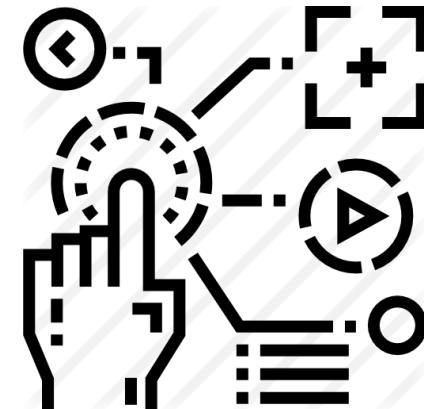


What Operating Systems Do

- Depends on the point of view
- Users want convenience, **ease of use** and **good performance**
 - Don't care about **resource utilization**
- But shared computer such as **mainframe** or **minicomputer** must keep all users happy
- Users of dedicate systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- Some computers have little or no user interface, such as embedded computers in devices and automobiles

Operating System Definition

- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer



What do we learn? Operating system



What is
OS?

Resource
allocator



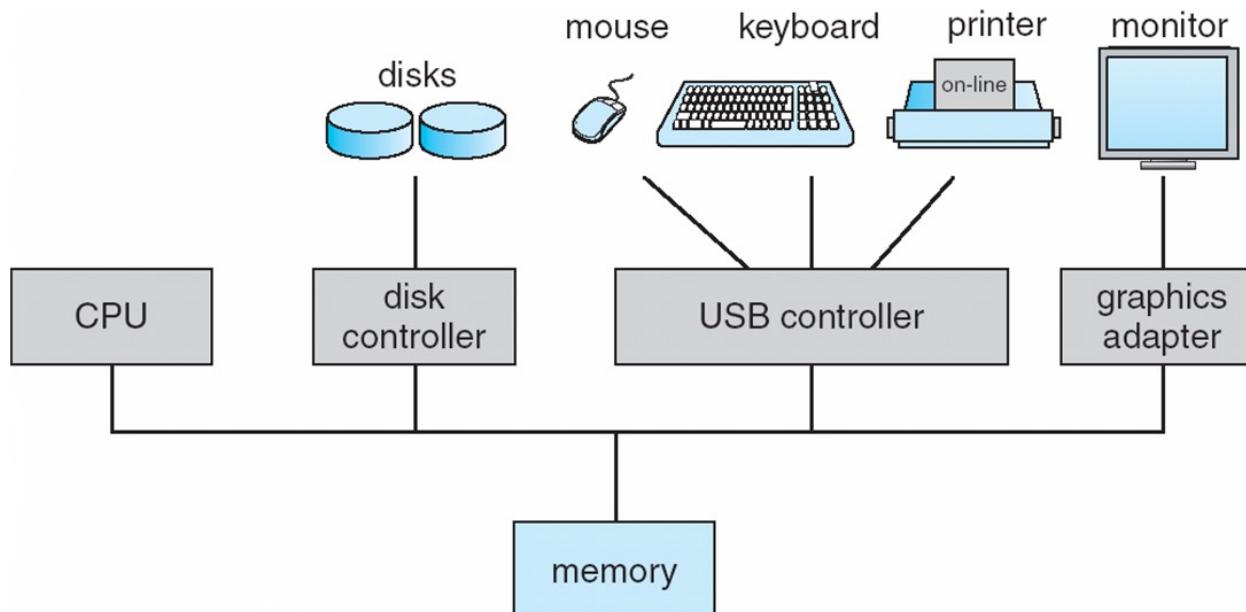
And
?

Control program

Computer System Organization

■ Computer-system operation

- One or more CPUs, device controllers connect through common bus providing access to shared memory
- Concurrent execution of CPUs and devices competing for memory cycles



Computer-System Operation

- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an **interrupt**

1.1.1 Why Interrupts?

- An operating system is **interrupt driven**
 - OS preserves state of CPU by storing registers and program counter
 - Determines which type of interrupt has occurred through polling of a vectored interrupt system
 - Different code performed action defined for different type of interrupt
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request



1.1.1 CPU – the brain

- **Central processing unit (CPU)**, also called a **central processor**, **main processor** or just **processor**,
- executes instructions that make up a computer program
- fetch, decode and execute steps in their operation, which are collectively known as the instruction cycle



1.1.2 Storage Structure

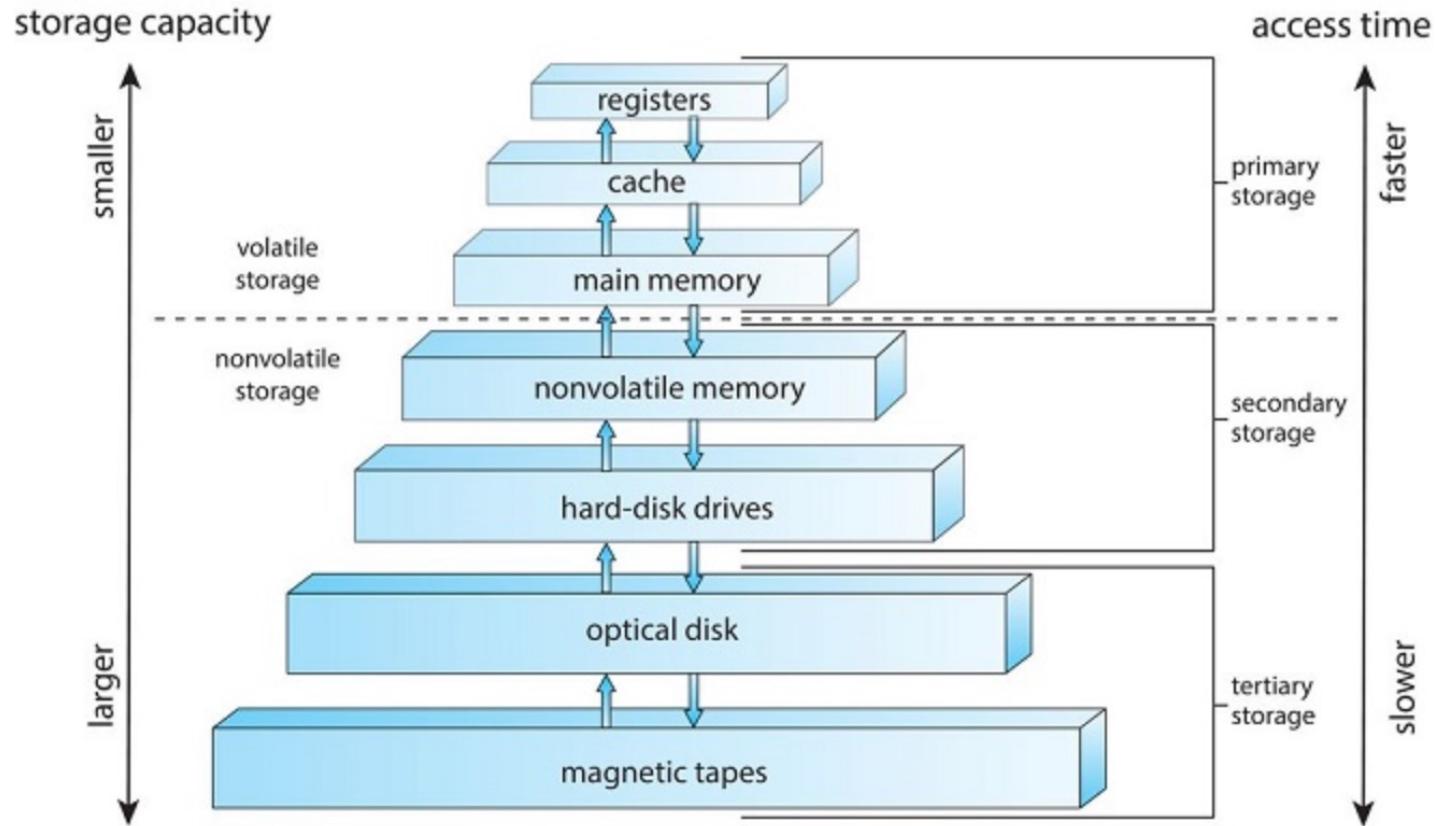
- Main memory – only large storage media that the CPU can access directly
 - **Random access memory (RAM)**
 - Typically **volatile**
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
 - Hard disk drive (HD) – rigid metal or glass platters covered with magnetic recording material
 - ▶ Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - **Solid-state disks** – faster than hard disks, nonvolatile (**We love it !!**)
 - ▶ Various technologies
 - ▶ Becoming more popular

Storage Hierarchy

- Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
- **Caching** – copying information into faster storage system; main memory can be viewed as a cache for secondary storage
- **Device Driver** for each device controller to manage I/O

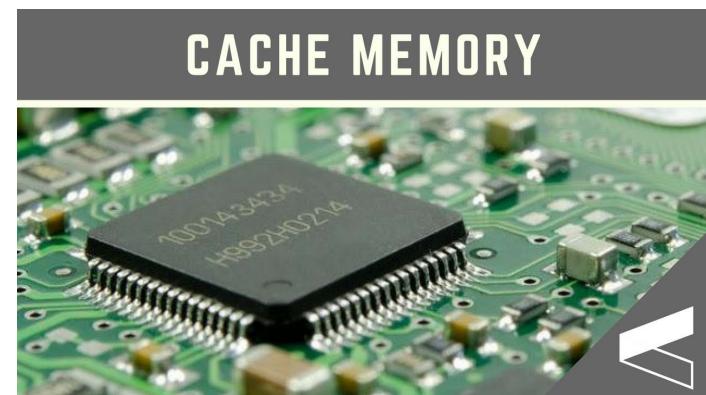


Storage-Device Hierarchy



Caching

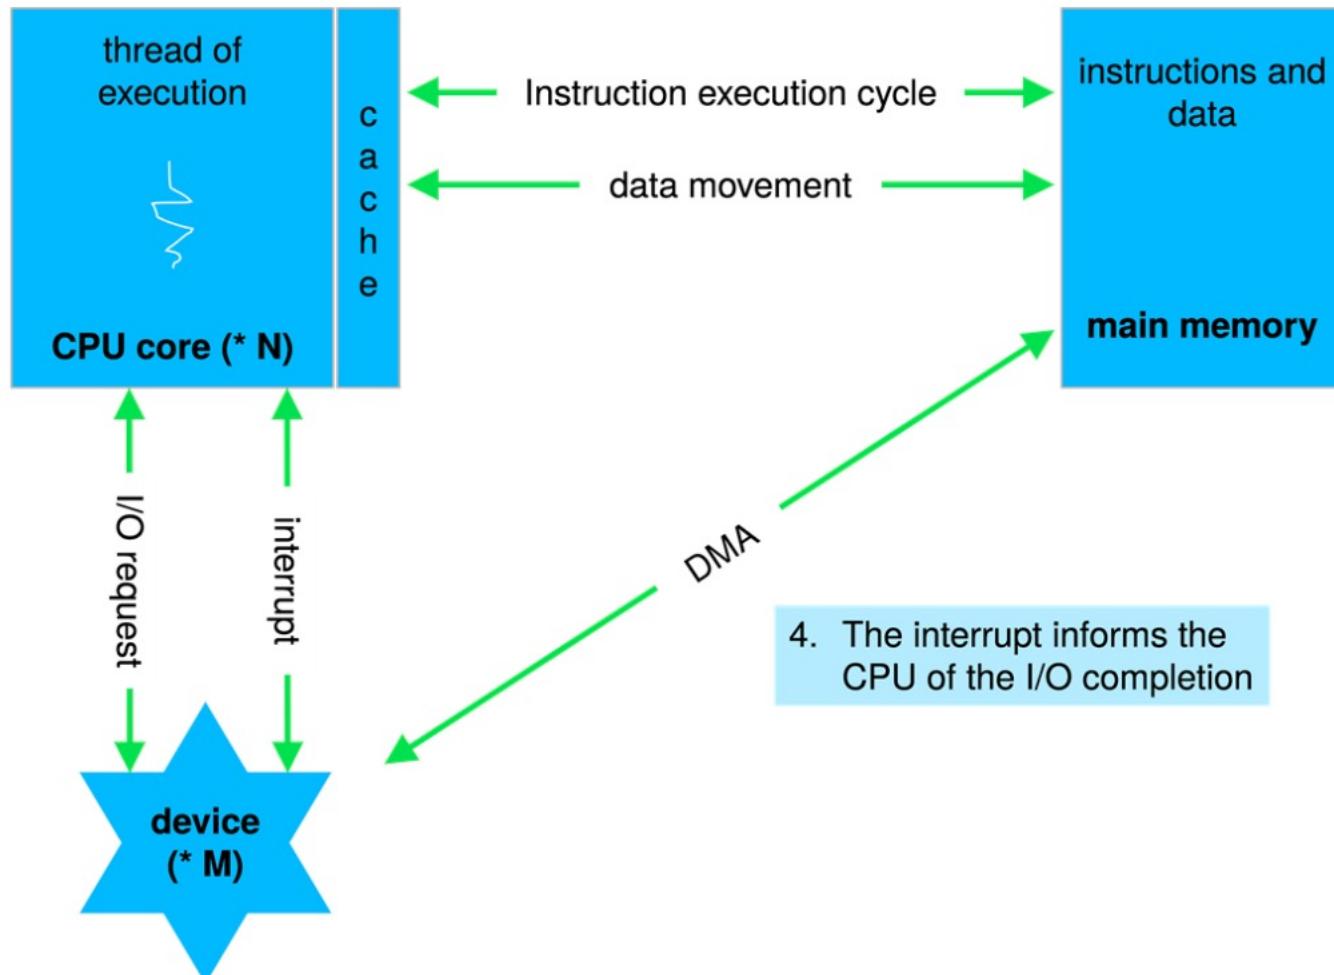
- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy



Direct Memory Access Structure

- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte

How a Modern Computer Works



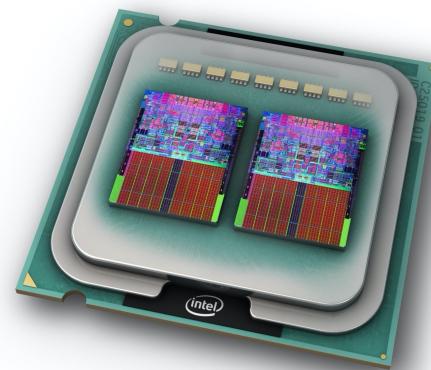
A von Neumann architecture

1.2 Computer-System Architecture

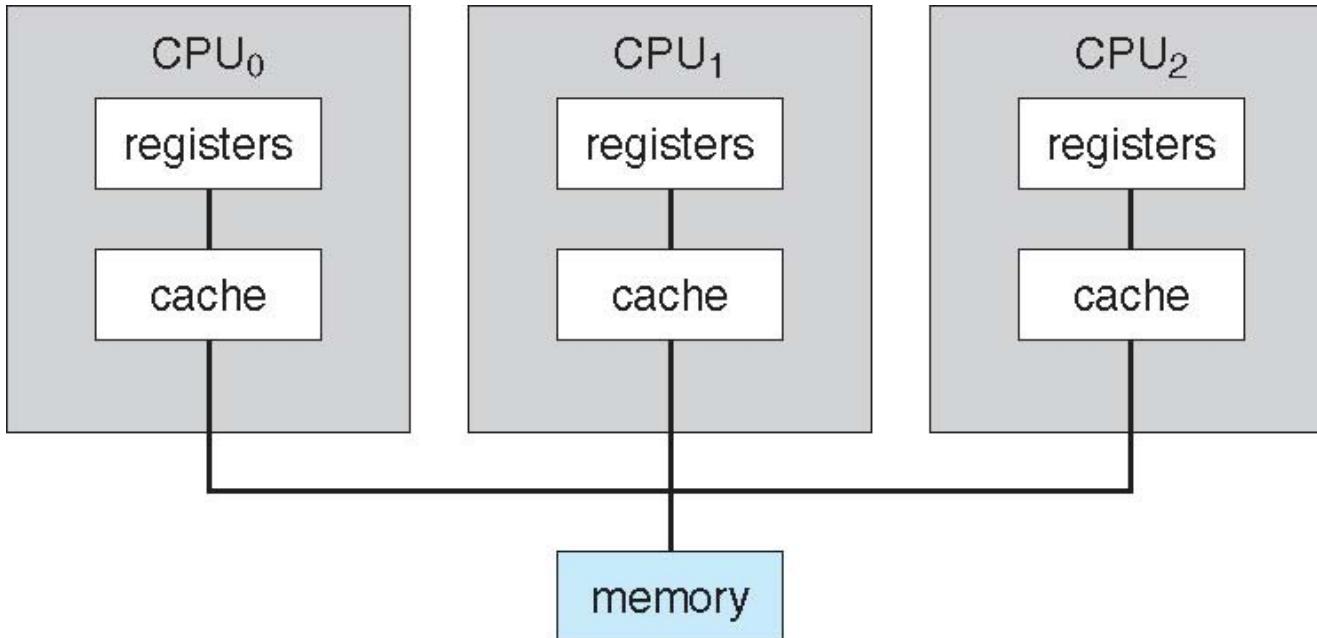
- **Single-processor system**

- **Multiprocessors system**

- Growing in use and importance
- Advantages include:
 1. Increased throughput
 2. Economy of scale
 3. Increased reliability – graceful degradation or fault tolerance



Computer-System Architecture



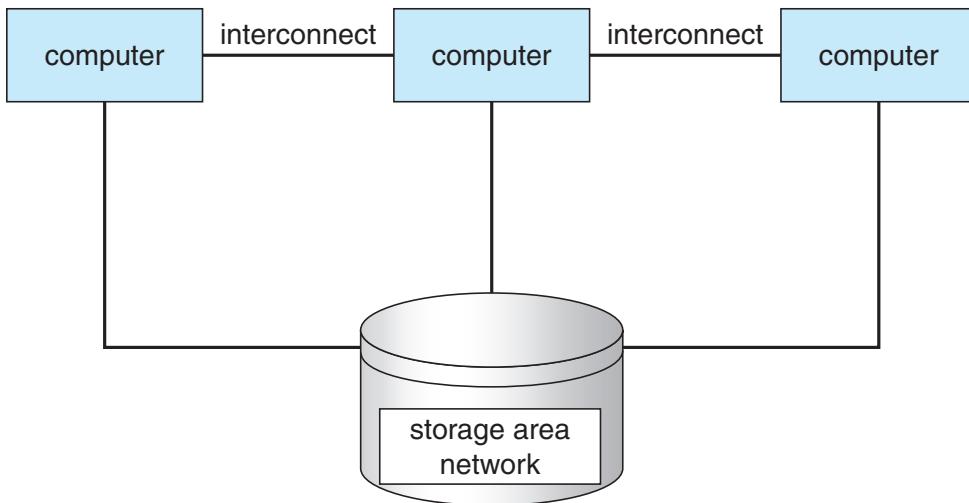
DEFINITIONS OF COMPUTER SYSTEM COMPONENTS

- **CPU**—The hardware that executes instructions.
- **Processor**—A physical chip that contains one or more CPUs.
- **Core**—The basic computation unit of the CPU.
- **Multicore**—Including multiple computing cores on the same CPU.
- **Multiprocessor**—Including multiple processors.

Although virtually all systems are now multicore, we use the general term **CPU** when referring to a single computational unit of a computer system and **core** as well as **multicore** when specifically referring to one or more cores on a CPU.

Computer-System Architecture

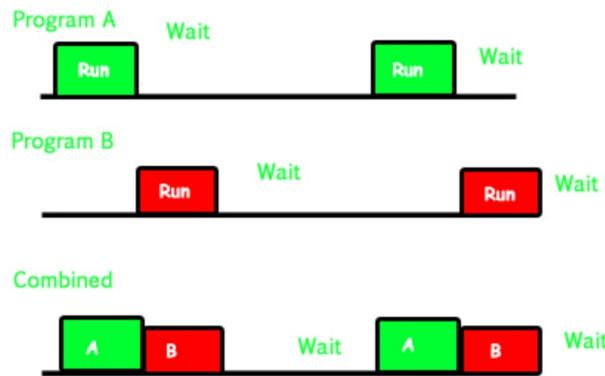
- **Clustered systems:** Like multiprocessor systems, but multiple systems working together
 - Usually sharing storage via a **storage-area network (SAN)**
 - Provides a **high-availability** service which survives failures
 - Some clusters are for **high-performance computing (HPC)**
 - ▶ Applications must be written to use **parallelization**



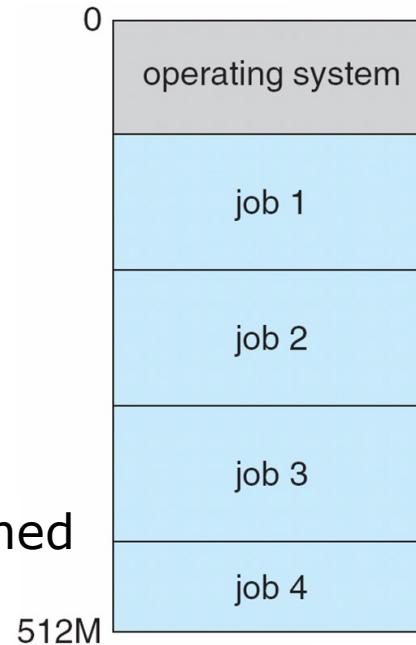
1.3 Operating System Structure

■ Multiprogramming (Batch system)

- Single user cannot keep CPU and I/O devices busy at all times
- Multiprogramming organizes jobs (code and data) so CPU always has one to execute
- A subset of total jobs in system is kept in memory
- One job selected and run via **job scheduling**
- When it has to wait (for I/O for example), OS switches to another job

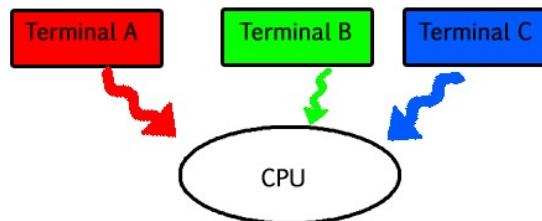


Memory Layout for Multiprogrammed System



1.3 Operating System Structure

- **Timesharing (multitasking)** is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive computing**
 - **Response time** should be < 1 second
 - Each user has at least one program executing in memory ⇒ **process**
 - If several jobs ready to run at the same time ⇒ **CPU scheduling**
 - If processes don't fit in memory, **swapping** moves them in and out to run
 - **Virtual memory** allows execution of processes not completely in memory



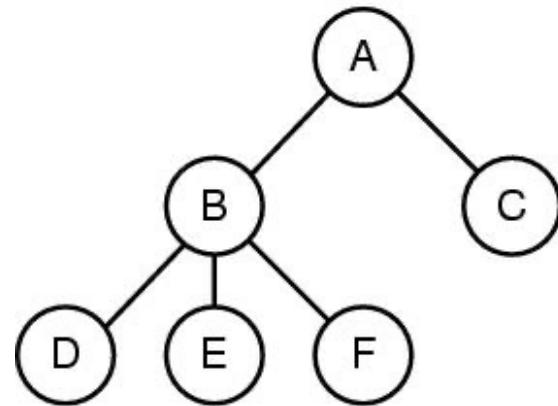
1.4 Resource Management

- Process Management
- Memory Management
- Storage Management
- I/O Management

1.4 Process Management

What is a process?

- Program in execution
- Lives in address space
- Process table
 - Keeps info about process
 - Used to re-start process
- Can communicate with one another



Process creates child processes
Tree structure

1.4 Process Management

- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread

Memory Management

- To execute a program all (or part) of the instructions must be in memory
- All (or part) of the data that is needed by the program must be in memory.
- Memory management determines what is in memory and when
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed

Storage Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit - **file**
 - Each medium is controlled by device (i.e., disk drive, tape drive)
 - ▶ Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
- File-System management
 - Files usually organized into directories
 - Access control on most systems to determine who can access what
 - OS activities include
 - ▶ Creating and deleting files and directories
 - ▶ Primitives to manipulate files and directories
 - ▶ Mapping files onto secondary storage
 - ▶ Backup files onto stable (non-volatile) storage media

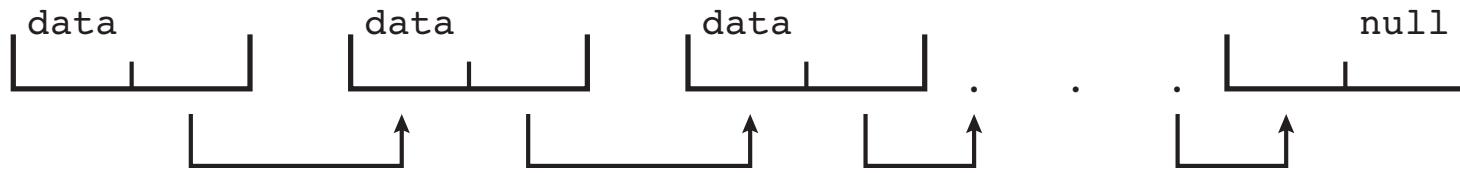
Performance of Various Levels of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

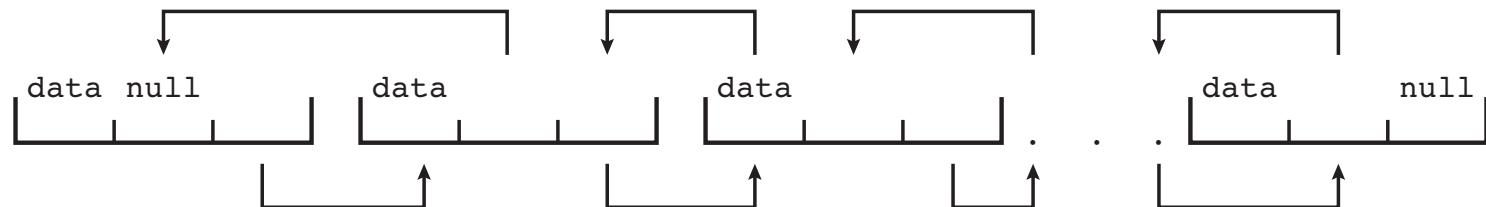
Movement between levels of storage hierarchy can be explicit or implicit

1.6 Kernel Data Structures

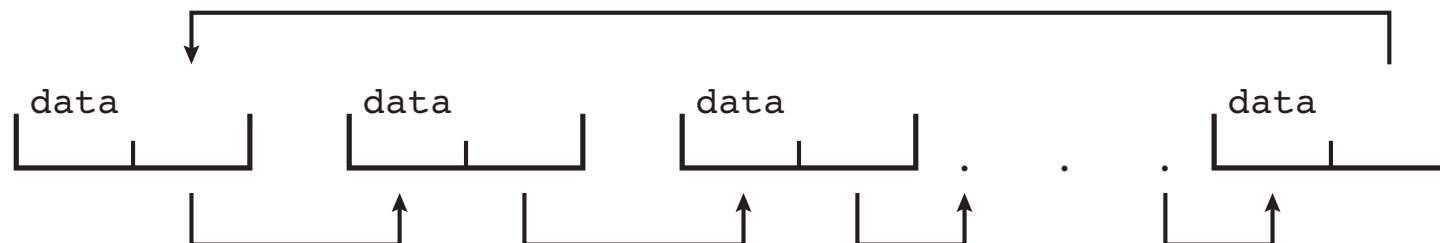
- Many similar to standard programming data structures
- *Singly linked list*



- *Doubly linked list*



- *Circular linked list*



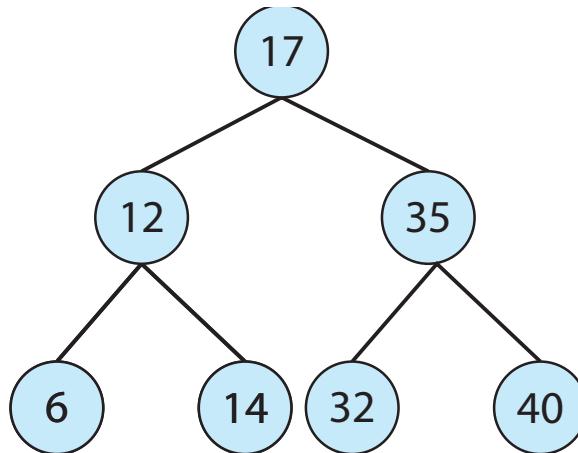
- *Stack and Queue*

1.6 Kernel Data Structures

■ Binary search tree

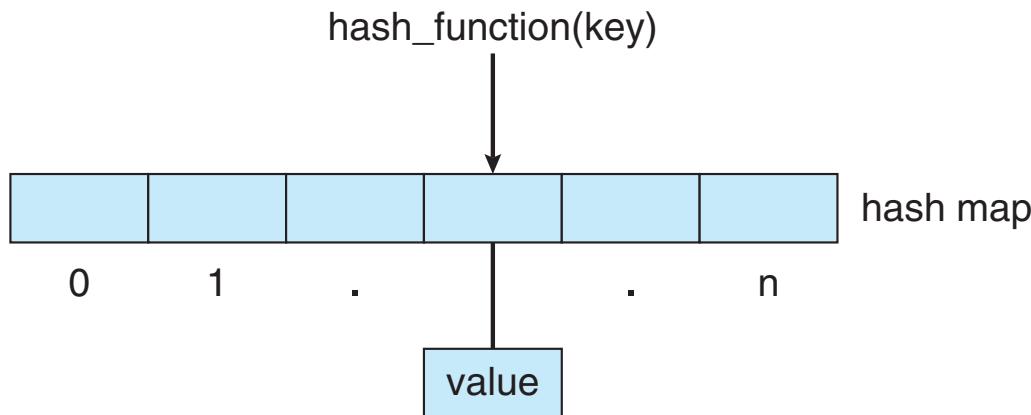
left \leq right

- Search performance is $O(n)$
- **Balanced binary search tree** is $O(\lg n)$



1.6 Kernel Data Structures

- Hash function can create a hash map



- Bitmap – string of n binary digits representing the status of n items
- Linux data structures defined in
 - include** files `<linux/list.h>`, `<linux/kfifo.h>`,
`<linux/rbtree.h>`

1.7 Computing Environments - Traditional

- Stand-alone general purpose machines
- But blurred as most systems interconnect with others (i.e., the Internet)
- Portals provide web access to internal systems
- **Network computers** are like Web terminals
- Mobile computers interconnect via **wireless networks**
- Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks



Computing Environments - Mobile

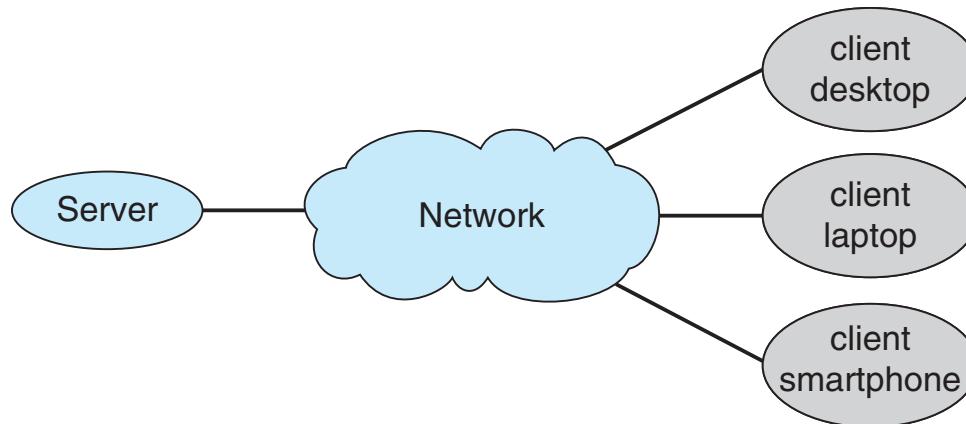
- Handheld smartphones, tablets, etc
- What is the functional difference between them and a “traditional” laptop?
- Extra feature – more OS features (GPS, gyroscope)
- Allows new types of apps like ***augmented reality***
- Use IEEE 802.11 wireless, or cellular data networks for connectivity
- Leaders are **Apple iOS** and **Google Android**



Computing Environments – Client-Server

■ Client-Server Computing

- Dumb terminals supplanted by smart PCs
- Many systems now **servers**, responding to requests generated by **clients**
 - ▶ **Compute-server system** provides an interface to client to request services (i.e., database)
 - ▶ **File-server system** provides interface for clients to store and retrieve files

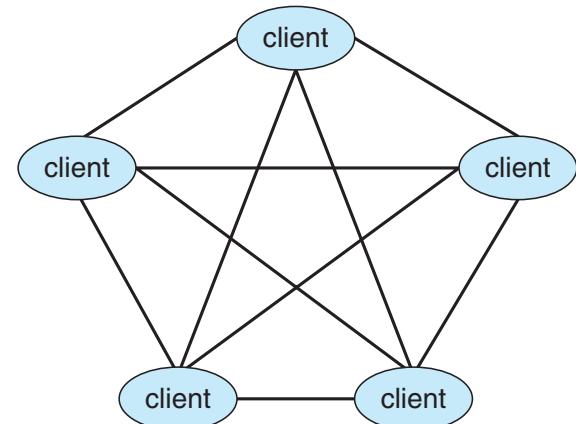


Computing Environments – Distributed

- Distributed computing
 - Collection of separate, possibly heterogeneous, systems networked together
 - ▶ **Network** is a communications path, **TCP/IP** most common
 - Local Area Network (LAN)
 - Wide Area Network (WAN)
 - Metropolitan Area Network (MAN)
 - Personal Area Network (PAN)
 - **Network Operating System** provides features between systems across network
 - ▶ Communication scheme allows systems to exchange messages
 - ▶ Illusion of a single system

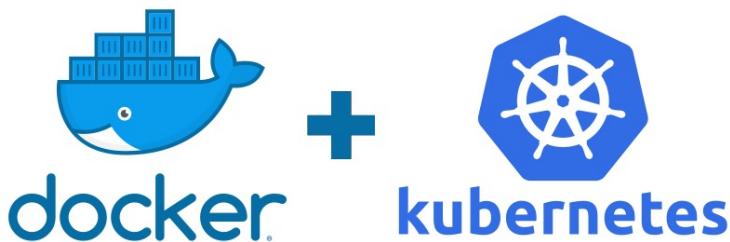
Computing Environments - Peer-to-Peer

- Another model of distributed system
- P2P does not distinguish clients and servers
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - ▶ Registers its service with central lookup service on network, or
 - ▶ Broadcast request for service and respond to requests for service via ***discovery protocol***
 - Examples include Bitcoin, Ethereum

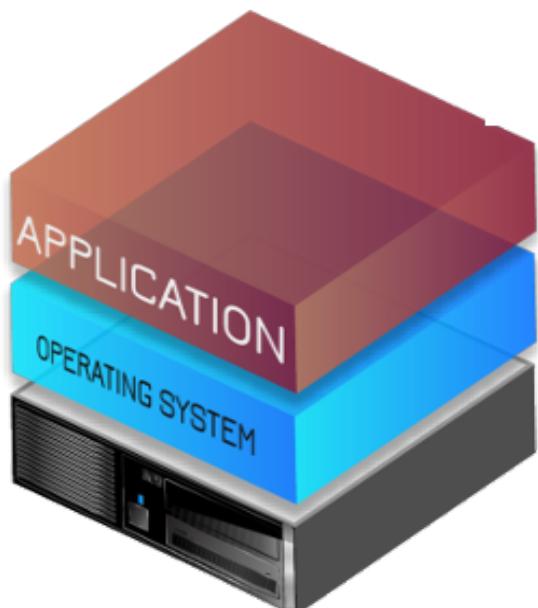


Computing Environments - Virtualization

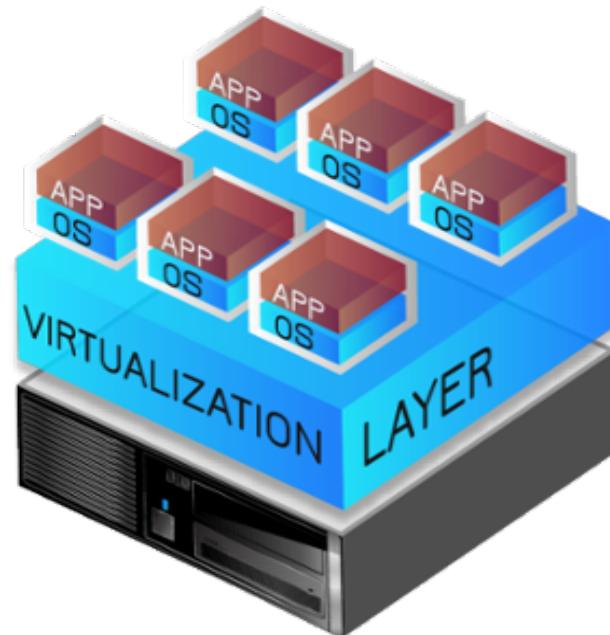
- Allows operating systems to run applications within other OSes
 - Vast and growing industry
- **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - **VMM** (virtual machine Manager) provides virtualization services



Computing Environments - Virtualization



**Traditional Server
Architecture**



**Virtualized Server
Architecture**

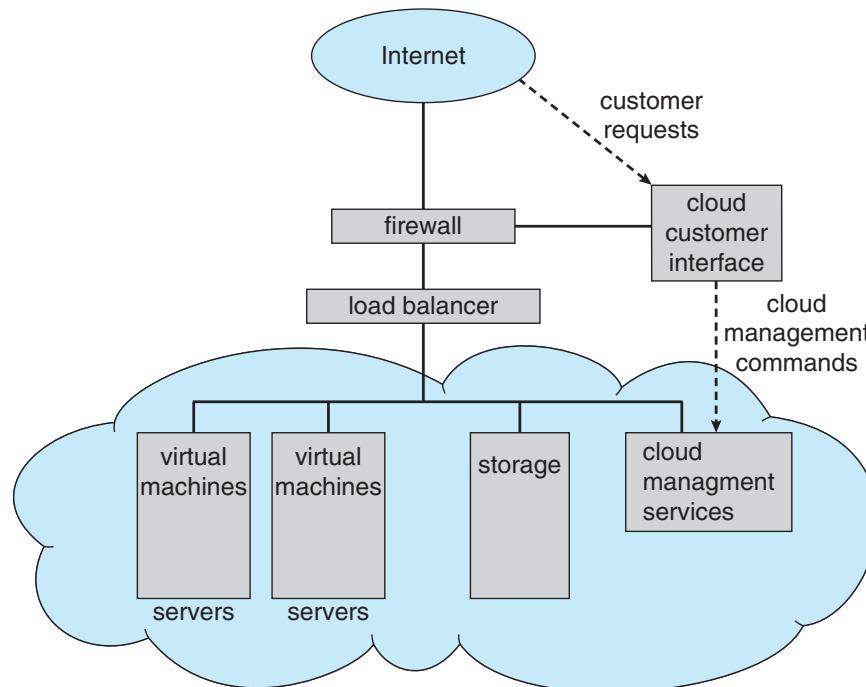
Computing Environments – Cloud Computing

- Delivers computing, storage, even apps as a service across a network
- Logical extension of virtualization because it uses virtualization as the base for its functionality.
 - Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage
- Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
 - Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)



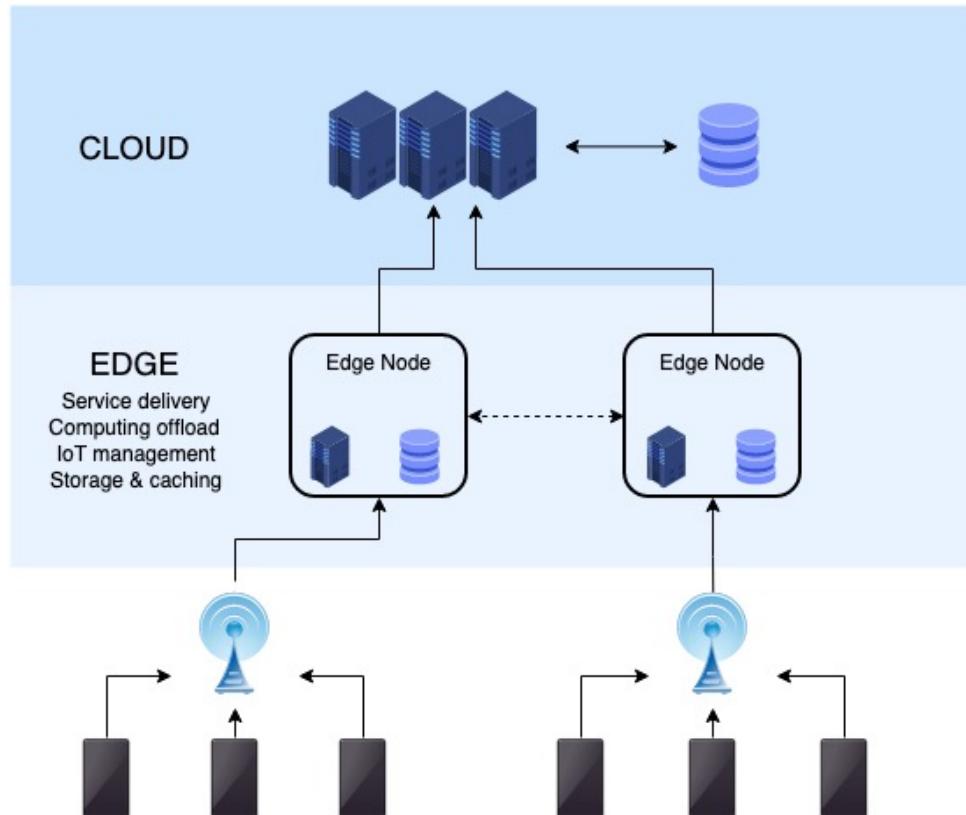
Computing Environments – Cloud Computing

- Cloud computing environments composed of traditional OSes, plus VMs, plus cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications



Computing Environments – Edge Computing / IoT

- a distributed computing paradigm that brings computation and data storage closer to the location where it is needed, to improve response times and save bandwidth
 - Through virtualization technology
 - run a wider range of applications on the edge servers



Open-Source Operating Systems

- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run guest operating systems for exploration



E.g., Apache Spark

- Spark and its RDDs were developed in 2012
 - Spark Core
 - RDDs
 - Spark SQL
 - Spark MLlib for machine learning
 - Spark Streaming



Matei Zaharia

Stanford DAWN Lab and Databricks

Verified email at cs.stanford.edu - [Homepage](#)

Distributed Systems Machine Learning Databases Security

[Spark: Cluster computing with working sets.](#)

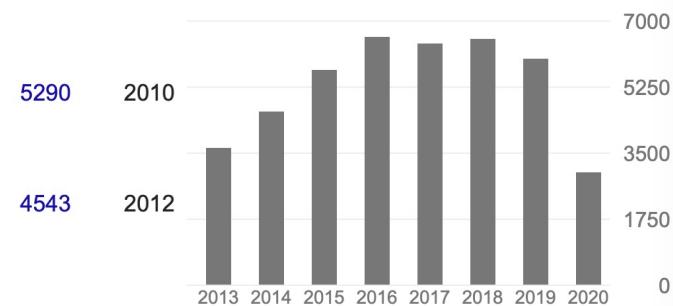
M Zaharia, M Chowdhury, MJ Franklin, S Shenker, I Stoica
HotCloud 10 (10-10), 95

[Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing](#)

M Zaharia, M Chowdhury, T Das, A Dave, J Ma, M McCauly, MJ Franklin, ...

Presented as part of the 9th {USENIX} Symposium on Networked Systems Design ...

Cited by	VIEW ALL	
	All	Since 2015
Citations	49440	34269
h-index	50	44
i10-index	95	87



1.49